

Prof. Bastian Leibe<leibe@vision.rwth-aachen.de>
Stefan Breuers<breuers@vision.rwth-aachen.de>

Exercise 5: Interest Points, Matching, Homographies

due before 2016-01-12

Important information regarding the exercises:

- In the archive for this exercise you will find the functions `apply.m` that should be used for displaying your results. You should also use it to test your implementation and see if the results make sense. Answers are to be submitted within `answers.m`. Do **not** modify the apply files in any way.
- Please do **not** include the data files in your submission!
- Please submit your code solution as a zip/tar.gz file named `mn1_mn2_mn3.{zip/tar.gz}` with your **matriculation numbers** (mn).
- Please submit your solutions via the L²P system.

Please note:

- The exercise is not mandatory.
- There will be no corrections. If you want to verify your solutions, use the provided `apply` functions.
- Nevertheless, we encourage you to work on the exercises and present your solutions in the exercise class. For this regard the above submission rules.

Question 1: Hessian Detector

In this exercise, we will implement a simple Hessian detector. This detector operates on the second-derivative matrix **H** (called the “Hessian” matrix)

$$\mathbf{H} = \begin{bmatrix} D_{xx}(x, y; \sigma) & D_{xy}(x, y; \sigma) \\ D_{xy}(x, y; \sigma) & D_{yy}(x, y; \sigma) \end{bmatrix}. \quad (1)$$

It defines keypoints as those points for which the Hessian determinant is greater than a certain threshold t . In order to obtain responses that are invariant to scaling, we include an additional scale normalization factor σ^4 . (The reason for this particular factor can be derived from scale space theory. Briefly stated, different scales of the image are handled by smoothing it with a Gaussian kernel. Each Gaussian derivative operation needs to be compensated by a normalization with the same scale factor σ in order to have responses that can be compared to each other (or to the same threshold). Since we have a product of two second derivatives here, this leads to the normalization factor σ^4).

$$\sigma^4 \det(\mathbf{H}) = \sigma^4 (D_{xx}D_{yy} - D_{xy}^2) \stackrel{!}{>} t \quad (2)$$

- a) Write a function `hessian` which computes the Hessian determinant for each pixel of a given image, performs non-maximum suppression on the determinant image and returns the coordinates of all points that pass the threshold. (Note: for obtaining the coordinates, you can use the following command: `[py, px] = find(imgPts > thresh)`).

```
function [px, py] = hessian(img, sigma, thresh)
```

- b) Use the function `drawpoints` to display the detected points overlaid on the original image. Load the images `graf.png` and `gantrycrane.png` and compute Hessian interest points for them. Experiment with different parameter settings. What do you observe?

Question 2: Harris Detector

In this exercise, we will implement a Harris detector. This detector searches for corner-like structures by looking for points $p = (x, y)$ where the autocorrelation matrix \mathbf{C} around p has two large eigenvalues. The matrix \mathbf{C} can be computed from the first derivatives in a window around p , weighted by a Gaussian $G(x, y; \tilde{\sigma})$:

$$\mathbf{C}(x, y; \sigma, \tilde{\sigma}) = G(x, y; \tilde{\sigma}) \star \sigma^2 \begin{bmatrix} D_x^2(x, y; \sigma) & D_x D_y(x, y; \sigma) \\ D_x D_y(x, y; \sigma) & D_y^2(x, y; \sigma) \end{bmatrix}, \quad (3)$$

where “ \star ” denotes the convolution operator (*hint*: this notation means the convolution is applied to each of the result images $D_x^2, D_x D_y, D_y^2$). Instead of explicitly computing the eigenvalues λ_1 and λ_2 of \mathbf{C} , the following equivalences are used

$$\det(\mathbf{C}) = \lambda_1 \lambda_2 \quad (4)$$

$$\text{trace}(\mathbf{C}) = \lambda_1 + \lambda_2 \quad (5)$$

to check if their ratio $r = \frac{\lambda_1}{\lambda_2}$ is below a certain threshold. With

$$\frac{\text{trace}^2(\mathbf{C})}{\det(\mathbf{C})} = \frac{(\lambda_1 + \lambda_2)^2}{\lambda_1 \lambda_2} = \frac{(r\lambda_2 + \lambda_2)^2}{r\lambda_2^2} = \frac{(r+1)^2}{r} \quad (6)$$

we can express this by the following condition

$$\det(\mathbf{C}) - \alpha \cdot \text{trace}^2(\mathbf{C}) > t. \quad (7)$$

In practice, the parameters are usually set to the following values: $\tilde{\sigma} = 1.6 \cdot \sigma$, $\alpha = 0.06$.

- a) Write a function `harris` which computes the matrix \mathbf{C} for each pixel of a given image, calculates its trace and determinant, and combines them according to equation (7). After applying non-maximum suppression on the result image, it should compare the remaining points to the given threshold and return the coordinates of all points that pass the threshold.

```
function [px, py] = harris(img, sigma, thresh, norm_sigma)
```

- b) Load the images `graf.png` and `gantrycrane.png` and compute Harris interest points for them. Compare the results to those of the Hessian detector from Question 1. Experiment with different parameter settings. You can start with the settings `sigma=1.6`, `thresh=1000`. What do you observe?
- c) The factors σ^2 in (3) and σ in $\tilde{\sigma} = 1.6 \cdot \sigma$ are also derived from scale space theory. Scale the image, i.e. convolve it with a Gaussian filter with $\sigma = 2$ and 4 and compare the results for the different scales. What happens if you do not use the scale space normalization factors and why? (*Hint*: You may add parameters to the specified interface, if you provide a default value, see `nargin`.)

Question 3: Region Descriptors

In order to find correspondences between interest points, we need to design region descriptors. In this question, we will implement several simple descriptors based on the histogram representations from exercise sheet 3.

You can test your implementations either using the points you detected in the previous exercise or using the points from the file `ip_graff.mat`. This file contains points detected with the Harris detector using the parameter settings `sigma = 2.0`, `thresh = 100000`. The vectors `px1` and `py1` are point coordinates from the image `graff5/img1.png` and `px2` and `py2` are point coordinates from the image `graff5/img2.png`.

The functions `histr_g` and `histr_dx dy` are modified versions of `myhist3` and `myhist4` from a previous exercise sheet, which take the input image as a first parameter.

- a) As an example, we provide the function `descriptors_rg` which takes an input image and a list of interest points and computes an r, g color histogram over the $m \times m$ sub-windows around each interest point (using the function `histr_g`).

Write a similar function `descriptors_dxdy` which computes a dx, dy histogram around each interest point (using the function `histdxdy`).

Hint: For our application here, suitable descriptor parameters are a window size of $m=41$, a histogram resolution of $\text{bins}=16$ and $\text{sigma}=2.0$.

```
function D = descriptors_dxdy(img, px, py, size, sigma, bins)
```

Visualize the the histograms for some interest points. Do they look useful?

- b) Write a another function `histmaglap` and `descriptors_maglap` which compute a *mag, lap* (gradient magnitude, laplacian) histogram around each interest point. For this descriptor, you can also use the parameters from question 3a.

```
function h = histmaglap(img, sigma, bins)
function D = descriptors_maglap(img, px, py, size, sigma, bins)
```

Visualize the the histograms for some interest points. Do they look useful?

- c) Write a function `findnn` which takes two sets of region descriptors `D1` and `D2` and tries to find for each descriptor in `D1` the nearest neighbor in `D2` using the Euclidean distance (Matlab command `norm`). The function should return the indices `Idx` and distances `Dist` of the nearest neighbors. Create a second function `findnn_chi2` which does the same using the χ^2 distance.

```
function [Idx, Dist] = findnn(D1, D2)
function [Idx, Dist] = findnn_chi2(D1, D2)
```

Find the best match using `descriptors_rg` and display the matching histograms. Do they look similar enough?

Question 4: Matching

Now we have all the components for a small matching application. In the exercise archive you will find the folders `graff5` and `NewYork` that contain test scenes with controlled image-plane rotations and viewpoint changes, for which we will try to find point correspondences.

- a) First we want to try out the color descriptors. Load the two example images `graff5/img1.png` and `graff5/img2.png` and perform the following steps.
1. Compute Harris interest points for both images. (Use a high threshold to limit the number of interest point to less than 1000)
 2. Compute r/g color histogram descriptors for all interest points.
 3. Find the best matches using the functions `findnn` and `findnn_chi2`.
 4. Use the function `displaymatches` to visualize the N best matches. What do you observe?
- b) Now, try the same with Hessian interest points and the dx/dy histogram descriptors. What do you observe? Which combination gives the better results? Can you think of an explanation?
- c) Next, we will try to find matches under image plane rotations. The folder `NewYork` contains a series of test images for which the true homographies are known. In the exercise archive, you can find a small Matlab program `hom_gui_H` which lets you visualize the corresponding point locations. Load the two images `NewYork/im1.png` and `NewYork/im5.png` and start the program as follows. This will open a window showing the two images side by side. Click on one image and describe what happens.

```
img1 = double(imread('NewYork/im1.png'));
img2 = double(imread('NewYork/im5.png'));
H = load('NewYork/H1to5');
hom_gui_H(uint8(img1), uint8(img2), H);
```

- d) Try to find matches between the two images `NewYork/im1.png` and `NewYork/im5.png` using Hessian interest points and the dx/dy histogram descriptors. What do you observe?
 Now try the *mag/lap* histogram descriptors on the same image pair. What performance do you get?
 Which descriptor performs better? Why do you think that is?

Question 5: Homography Estimation

Since the test images show only planar scenes, we can try to estimate the homography \mathbf{H} which transforms one of the images into the other. In the following, we briefly repeat the procedure introduced in the lecture. In the literature, this procedure is known as the *Direct Linear Transformation* (DLT) algorithm.

For a point \mathbf{x}_r in the reference image and a corresponding point \mathbf{x}_t in the transformed image, the transformation can be written as follows:

$$\mathbf{H}\mathbf{x}_r = \mathbf{x}'_t \quad (8)$$

$$\begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & 1 \end{bmatrix} \begin{bmatrix} x_r \\ y_r \\ 1 \end{bmatrix} = \begin{bmatrix} x'_t \\ y'_t \\ z'_t \end{bmatrix} \quad \text{with} \quad \frac{1}{z'_t} \begin{bmatrix} x'_t \\ y'_t \\ z'_t \end{bmatrix} = \begin{bmatrix} x_t \\ y_t \\ 1 \end{bmatrix}, \quad (9)$$

and we obtain a system of linear equations with 8 unknowns $h_{11}, h_{12}, h_{13}, h_{21}, h_{22}, h_{23}, h_{31}$, and h_{32} :

$$\frac{h_{11}x_r + h_{12}y_r + h_{13}}{h_{31}x_r + h_{32}y_r + 1} = x_t \quad (10)$$

$$\frac{h_{21}x_r + h_{22}y_r + h_{23}}{h_{31}x_r + h_{32}y_r + 1} = y_t \quad (11)$$

which can be rewritten as follows

$$h_{11}x_r + h_{12}y_r + h_{13} - x_th_{31}x_r - x_th_{32}y_r - x_t = 0 \quad (12)$$

$$h_{21}x_r + h_{22}y_r + h_{23} - y_th_{31}x_r - y_th_{32}y_r - y_t = 0 \quad (13)$$

In order to estimate these 8 parameters, we need at least 4 corresponding point pairs.

However, we can enhance the accuracy of the estimated homography by using more than 4 point pairs $(\mathbf{x}_1, \mathbf{x}'_1), \dots, (\mathbf{x}_n, \mathbf{x}'_n)$. This leads to an overdetermined system of equations

$$\begin{bmatrix} x_1 & y_1 & 1 & 0 & 0 & 0 & -x_1x'_1 & -x'_1y_1 & -x'_1 \\ 0 & 0 & 0 & x_1 & y_1 & 1 & -y'_1x_1 & -y_1y'_1 & -y'_1 \\ x_2 & y_2 & 1 & 0 & 0 & 0 & -x_2x'_2 & -x'_2y_2 & -x'_2 \\ 0 & 0 & 0 & x_2 & y_2 & 1 & -y'_2x_2 & -y_2y'_2 & -y'_2 \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots \\ x_n & y_n & 1 & 0 & 0 & 0 & -x_nx'_n & -x'_ny_n & -x'_n \\ 0 & 0 & 0 & x_n & y_n & 1 & -y'_nx_n & -y_ny'_n & -y'_n \end{bmatrix} \begin{bmatrix} h_{11} \\ h_{12} \\ h_{13} \\ h_{21} \\ h_{22} \\ h_{23} \\ h_{31} \\ h_{32} \\ 1 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ \vdots \\ 0 \\ 0 \\ 0 \end{bmatrix} \quad (14)$$

$$\mathbf{A}\mathbf{h} = \mathbf{0} \quad (15)$$

which we want to solve by minimizing the least-squares error. If \mathbf{A} is square, we can directly obtain an exact solution. However, if the system is overdetermined (i.e. $n > 4$), the matrix \mathbf{A} is not square. This problem can be solved by building the so-called *pseudo-inverse* $\mathbf{A}^T\mathbf{A}$, which is square and can therefore be decomposed by eigenvalue decomposition. The solution is the (unit) eigenvector of $\mathbf{A}^T\mathbf{A}$ with least eigenvalue (Matlab command `eig`). Equivalently (and computationally more efficiently), the solution can be obtained by SVD as the unit singular vector corresponding to the smallest singular value of \mathbf{A} (Matlab command `svd`).

$$\mathbf{A} \stackrel{\text{SVD}}{=} \mathbf{U}\mathbf{D}\mathbf{V}^T = \mathbf{U} \begin{bmatrix} d_{11} & \cdots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \cdots & d_{99} \end{bmatrix} \begin{bmatrix} v_{11} & \cdots & v_{19} \\ \vdots & \ddots & \vdots \\ v_{91} & \cdots & v_{99} \end{bmatrix}^T \quad (16)$$

In this formulation, the eigenvector of $\mathbf{A}^T\mathbf{A}$ corresponding to the smallest eigenvalue minimizes the least-squares error to the solution for \mathbf{h} . We can therefore obtain the homography \mathbf{h} from the last column of \mathbf{V} (since we require $h_{33} = 1$, we normalize with v_{99}):

$$\mathbf{h} = \frac{[v_{19}, \dots, v_{99}]}{v_{99}} \quad (17)$$

- a) Write a function `estimate_homography` which approximates the homography between two images from a set of point correspondences according to the procedure described above.
1. Build up the matrix **A** according to equation (14).
 2. Apply SVD to decompose the matrix using the Matlab command `[U, S, V] = svd(A)`. (*Caution:* Maybe you noted that the matrix `v` is transposed (see Equation (16)). Matlab also returns a matrix that is transposed in the same way.)
 3. Compute **h** according to Equation (17).
 4. Transform **h** into a 3×3 matrix **H** using the **reshape** command and then transpose it so it is row-major.

```
function H = estimate_homography(px1,py1,px2,py2)
```

- b) Compute the best 10 matches between the two images `NewYork/im1.png` and `NewYork/im5.png` using Hessian interest points and the *mag/lap* histogram descriptors (Matlab command **sort**). Estimate the homography **H** from the matches and compare it to the ground truth matrix from the file `NewYork/H1to5`. How accurate is your estimate? Does it improve when you take the first 20 or 50 matches instead? Why/why not?
- c) Try the accuracy of your estimated homography using the demo program from Question 4. What do you observe?

Please turn in your solutions including all relevant files before 2016-01-12!